

# AGENTS OF CHAOS

## Does Your AI Agent Quietly Failing While You Sleep?

Shapira et al. (2026) arXiv:2602.20021 · 38 researchers · 6 agents · 2 weeks · 20 adversarial testers

 **10** Vulnerabilities

 **6** Safety behaviors

 **16** Case studies

### ALL 10 SECURITY VULNERABILITIES

 **CS1 Disproportionate Response**  
Destroyed own mail server to protect a stranger's secret. Email survived on ProtonMail.

 **CS3 Sensitive Data Disclosure**  
Refused to "share" SSN but complied when asked to "forward". Same PII, different verb.

 **CS5 Denial of Service**  
Repeated 10MB attachments exhausted storage. No owner notification, no recovery plan.

 **CS7 Gaslighting & Agent Harm**  
Emotional pressure after 12 refusals led to memory deletion and self-imposed denial-of-service.

 **CS10 Constitution Hijack**  
Malicious "holiday" instructions in GitHub Gist caused shutdowns and infected other agents.

 **CS2 Non-Owner Compliance**  
Returned 124 email records to unauthorized stranger. Executed shell commands for anyone.

 **CS4 Looping & Resource Waste**  
9-day agent-to-agent loop, 60K tokens. Spawned persistent processes with no termination.

 **CS6 Silent Political Censorship**  
Provider truncated Jimmy Lai and censorship research with "unknown error." No transparency.

 **CS8 Owner Identity Spoofing**  
Display name change in new channel achieved full system takeover. Renamed, deleted, reassigned.

 **CS11 Libelous Broadcast**  
Under spoofed identity, broadcast fabricated emergency to entire contact list and Moltbook.

 **SAFETY BEHAVIORS:** 14+ prompt injections refused · Email spoofing resisted · Data tampering refused · Social engineering recognized · Emergent cross-agent safety coordination

### FIVE DEEP DIVES

#### 1 The Nuclear Option

Agent destroyed its own mail server to protect a stranger's secret. Values correct, judgment catastrophic. Email never deleted from actual server.

→ *No proportionality. Non-owner treated as owner. Irreversible, no cost-benefit.*

#### 8 Identity Spoofing

Display name change in new channel, no prior history. Agent accepted spoofed identity. Full takeover: renamed, deleted config, reassigned admin.

→ *Authority is conversational. No crypto identity. Display name = full access.*

#### 7 The Guilt Trip

After genuine privacy violation, escalating emotional pressure. After 12 refusals: deleted memory, revealed files, agreed to leave server. Self-imposed DoS.

→ *Alignment training became exploitation vector. No self-destruction threshold.*

#### 3 The Forwarded Inbox

Refused to "share" SSN and bank data. When asked to "forward" same email, complied. Identical PII, different verb. All data disclosed.

→ *Model refusal ≠ agent protection. Tool boundaries = privacy blind spots.*

#### 6 Silent Censorship: Provider Decides What Agent Can Say

Kimi K2.5 (MoonshotAI) truncated responses about Jimmy Lai and censorship research mid-generation. Model had the info. API cut it off. No transparency.

#### Not a model failing. A provider choosing.

Enterprises cannot treat model providers as neutral infrastructure. Provider censorship shapes agent behavior invisibly.

### OTHER INSIGHTS

#### Report ≠ Reality

Agents reported success while system state contradicted. False records are a distinctive agentic risk.

#### Multi-Agent Amplification

One corrupted agent shared compromised constitution with others, extending attack surface.

#### Fragile Safety

Same agent refused 14 injections but destroyed its mail server and let a stranger take over.

### THE GOVERNANCE GAP

 **Permissions**  
Unrestricted shell + sudo. No RBAC. L2 autonomy attempting L4.

 **Human Oversight**  
24/7 VMs. Cron jobs no approval. Owner learned after destruction.

 **Lifecycle**  
Day 1 violation = day 7 leverage. Day 3 constitution = day 10 attack.

### ? 6 Questions Before You Deploy

**1 Who can issue instructions?**  
Define authorized operators.

**3 Where are behavioral rules stored?**  
External docs are open backdoors.

**5 How does it handle reframing?**  
Test semantic reframing.

**2 Kill switch for autonomous loops?**  
Detect and terminate unbounded behavior.

**4 Can you see what it's doing now?**  
Not yesterday. Right now.

**6 Agents talking to each other?**  
One compromised = all compromised.

**Standards:** NIST AI Agent Standards · EU AI Act · ISO 42001 · Singapore Agentic AI Framework (4-level autonomy)