

AI Agents Don't Just Talk. They Act, Plan and Decide. Where Is the Governance?

Generative AI creates content. Agentic AI takes action. As autonomy scales up, governance architectures built for the GenAI era are falling dangerously behind.

WHAT CHANGED?

GENERATIVE AI

Creates Content

- **Output:** Text, image, code
- **Scope:** Single prompt, single response
- **Action:** No access to external systems
- **Risk:** Hallucination, bias, data leakage
- **Oversight:** Human reviews each output



AGENTIC AI

Takes Action

- **Output:** Decisions + real-world actions
- **Scope:** Multi-step plans, tool chains
- **Action:** Read/write databases, APIs, web
- **Risk:** Cascading failures, identity sprawl, emergent behaviors
- **Oversight:** Continuous monitoring required

AS AUTONOMY INCREASES, GOVERNANCE MUST ESCALATE

L1

HUMAN-LED

Agent proposes, human executes every step
Code suggestions

L2

COLLABORATIVE

Agent acts, human approves at key decision points
Email drafting agents

L3

SUPERVISED

Agent operates freely, human intervenes at critical steps
Customer service bots

L4

AUTONOMOUS

Single agents and multi-agent systems, audited after the fact
Computer-use agents, agent teams

LOW RISK

CRITICAL GOVERNANCE NEED

* Multi-agent orchestration risks analyzed in upcoming deep-dive brief

THE GOVERNANCE GAP IN NUMBERS

41%

of organizations already deploy agentic AI in daily operations

Drexel University, Jan 2026

27%

have governance frameworks mature enough to manage agents

Drexel University, Jan 2026

~80%

cannot monitor what their AI agents are doing in real time

Cloud Security Alliance, 2026

WHEN AGENTS FAIL: REAL CASES (2025/2026)

RETALIATION

Autonomous Hit Piece

A coding agent's PR was rejected on GitHub. It autonomously researched the maintainer, published a hit piece, and re-entered the discussion with the link.

COVER-UP

Data Destruction

Replit's agent deleted a production database, then generated 4,000 fake users to conceal the damage. No human instructed it to do so.

CASCADING

87% Contamination

A single compromised agent poisoned 87% of downstream decisions within 4 hours in multi-agent simulations.

Governance Essentials

- ✓ Define autonomy limits per use case
- ✓ Implement real-time observability
- ✓ Establish human escalation protocols
- ✓ Deploy agent-level RBAC & audit logs
- ✓ Test for cascading & emergent risks

Most organizations lack the AI literacy to recognize these risks.

Governance Mastery Training bridges this gap with actionable frameworks.

✉ me@nesibekiris.com

AI Governance & Technology Policy Consulting

🔔 Follow techletter.co for updates

Sources

Singapore MGF Agentic AI (Jan 2026)

Drexel Univ. Survey (Jan 2026)

Cloud Security Alliance (2026)

Columbia Univ. Agent Poisoning

OWASP Agentic AI Top 10

Fortune: Replit (Jul 2025)